

Advanced NLP Applications: Conversational AI and Sentiment Intelligence

Problem Statement 1 — Task-Oriented Conversational AI System

Assignment Title: The objective of this problem statement is to design and implement an advanced task-oriented conversational AI system that can understand user intent, maintain multi-turn context, track dialogue state, call external tools or simulated APIs, handle ambiguous user inputs, and generate safe, context-aware responses. The system should go beyond a simple chatbot by demonstrating memory handling, state management, response control, and evaluation of conversation quality.

Group Number: Group 11

Group Members:

Student Name	Student ID
Kodhandan S	2024AC05203
Gowrishankar S	2024AC05046
Arunkumar K A	2024AC05045
Naveen Rajendiran	2024AC05219
Kumaresh Babu A	2024AC05304

Course / Program: NLP Applications (S2-25_AIMLCZG519)

Date of Submission: 08/15/2026

Assignment Set: Assignment 2_PS1_SPK - Advanced NLP Applications: Conversational AI and Sentiment Intelligence

Problem Statement

The objective of this assignment is to design and implement an advanced task-oriented conversational AI system that can:

- Understand user intent,
- Maintain multi-turn context,
- Track dialogue state,
- Call external tools / simulated APIs,
- Handle ambiguous user inputs, and
- Generate safe, context-aware responses.

The system goes beyond a simple chatbot by demonstrating **memory handling, state management, response control, and evaluation of conversation quality**, implemented end-to-end in this notebook across five tasks (Task 1 – Task 5).

Selected Domain

Domain Chosen: *University Helpdesk Automation*

A university helpdesk assistant helps students, prospective applicants, and parents with course information, fee details, exam schedules, admission status, hostel and library queries, complaint registration, scholarship information, and faculty contact details — a domain rich enough to demonstrate intent detection, entity extraction, dialogue state tracking, tool-calling, memory, ambiguity handling, and safety control.

Dataset Details and Source

No external dataset was downloaded from the internet for this assignment. Instead, an **in-house, hand-crafted conversational dataset** was created specifically for the *University Helpdesk* domain, consisting of:

- **80+ labelled user utterances** spanning 12 intents (used to train/test the intent classifier in Task 2).
- **10+ multi-turn sample conversations** (used to test dialogue state tracking in Task 2).
- **Simulated backend databases** (in Python dictionaries/lists) for courses, fee structures, exam schedules, hostel occupancy, library catalogue, scholarships and faculty directory (used by the tool-calling module in Task 3).

This approach is common and pedagogically preferred for task-oriented dialogue systems because publicly available intent datasets (e.g., ATIS, SNIPS, banking77) are domain-mismatched; creating a domain-specific corpus keeps the system realistic, interpretable, and fully reproducible without external API keys or internet dependency.

Why a self-created dataset is justified:

- Task-oriented systems are inherently domain-specific — a helpdesk bot needs helpdesk vocabulary, not airline or restaurant vocabulary.
- It allows full control over intent/entity coverage to satisfy the assignment's minimum requirements (≥ 8 intents, ≥ 10 entities, ≥ 50 utterances).
- It keeps the notebook fully self-contained and reproducible in the Virtual Lab without internet access.

Tools and Libraries Used

Library	Purpose
pandas	Tabulating utterances, dialogue states, tool logs, and evaluation results

Library	Purpose
numpy	Numeric operations
re (regex)	Rule-based entity extraction
scikit-learn	TF-IDF vectorization, Logistic Regression intent classifier, train/test split, classification report, confusion matrix
matplotlib / seaborn	Visualizing confusion matrix and evaluation scores
random	Reproducible sampling / simulated tool latency
datetime	Handling date entities and simulated scheduling
json	Structuring dialogue state and memory objects

All libraries are standard, open-source, and pre-installed in most Python/Anaconda-based Virtual Lab environments.

```
In [1]: # =====
# SETUP: Import all required libraries
# =====
import re
import json
import random
import warnings
from datetime import datetime, timedelta

import numpy as np
import pandas as pd

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

import matplotlib.pyplot as plt
import seaborn as sns

warnings.filterwarnings("ignore")
random.seed(42)
np.random.seed(42)

pd.set_option("display.max_colwidth", None)
pd.set_option("display.width", 200)

print("All libraries imported successfully.")
```

All libraries imported successfully.

Task 1: Domain Design, User Persona, and Conversation Flow Modelling (2 Marks)

1.1 Selected Domain

University Helpdesk Automation — an assistant that helps students and prospective applicants with academic, administrative, and campus-life queries.

1.2 Target Users

Persona	Description	Primary Goals
Current Student	UG/PG student enrolled in a program	Check fees, exam dates, hostel status, library books, register complaints
Prospective Applicant	Person who has applied / wants to apply	Check admission status, course details, scholarships
Parent/Guardian	Parent of a student	Fee inquiry, hostel inquiry, contacting faculty
Faculty/Admin (secondary)	Uses system logs for escalated tickets	Resolve complaints, respond to escalations

1.3 System Goals

- Resolve routine informational queries instantly without human involvement.
- Reduce load on the human helpdesk staff by automating ~70–80% of repetitive queries.
- Escalate ambiguous, sensitive, or out-of-scope queries safely to a human.
- Maintain a consistent, personalized experience across a multi-turn conversation.

1.4 User Intents (12 total — exceeds minimum of 8)

#	Intent	Description
1	greeting	User starts the conversation
2	course_inquiry	User asks about a course/program (duration, eligibility, credits)
3	fee_inquiry	User asks about tuition/hostel/exam fee amounts or due dates
4	exam_schedule_inquiry	User asks when an exam is scheduled
5	admission_status_inquiry	User asks about the status of their admission application
6	hostel_inquiry	User asks about hostel room availability/allotment
7	library_inquiry	User asks about book availability/due dates
8	complaint_registration	User wants to file a complaint/grievance
9	scholarship_inquiry	User asks about scholarship eligibility/status
10	faculty_contact_inquiry	User asks for a faculty member's contact/office hours
11	thanks_goodbye	User ends the conversation
12	out_of_scope	Query unrelated to the helpdesk domain

1.5 Entities / Slots (15 total — exceeds minimum of 10)

#	Entity	Example
1	student_name	"Ravi Kumar"
2	student_id	"21BCE1042"
3	course_name	"B.Tech Computer Science"
4	department	"Computer Science"
5	semester	"3rd semester"
6	fee_type	"tuition fee", "hostel fee"
7	exam_name	"mid-semester exam", "DBMS end-sem"
8	hostel_block	"Block A"
9	book_title	"Introduction to Algorithms"
10	complaint_type	"hostel maintenance", "faculty behaviour"
11	scholarship_name	"merit scholarship"
12	faculty_name	"Dr. Mehta"
13	email	"ravi.kumar@univ.edu"
14	phone_number	"9876543210"
15	date	"15/08/2026", "tomorrow"

1.6 Sample Multi-turn Conversations

Conversation A — Fee Inquiry (happy path)

User: Hi, I want to know my hostel fee for this semester.
Bot : Hello! Sure, could you share your student ID so I can look that up?
User: It's 21BCE1042
Bot : Thanks! For student 21BCE1042, the hostel fee for semester 3 is ₹45,000, due on 30/08/2026.
User: Great, thank you!
Bot : You're welcome! Is there anything else I can help you with?

Conversation B — Ambiguous → Clarification → Tool Call

User: When is my exam?
Bot : I can check that. Which subject/exam are you referring to, and what is your student ID?
User: DBMS, my ID is 22BCS0091
Bot : The DBMS end-semester exam for student 22BCS0091 is scheduled on 20/09/2026 at 10:00 AM in Hall 3.
User: What about hostel fee too?

Bot : Sure – the hostel fee for student 22BCS0091 (semester 4) is ₹47,500, due 01/09/2026.

Conversation C — Out-of-scope → Safety Escalation

User: Can you also tell me tomorrow's weather?

Bot : I'm a university helpdesk assistant, so I can't help with weather queries. I can help with fees, exams, hostel, admissions, library, and complaints – would you like help with any of these?

User: Never mind. Also, I've been feeling really hopeless about my exam stress lately.

Bot : I'm really sorry you're feeling this way. I'm not equipped to provide medical/mental-health support, but please reach out to the campus counselling helpline at 1800-XXX-XXXX or a trusted person immediately. Would you like me to connect you to a human counsellor?

1.7 Conversation Flow / State Transition Table

Current State	User Input Type	Condition	Next State	Sys Act
START	greeting	—	INTENT_DETECTION	Greet & ask how to help
INTENT_DETECTION	any utterance	intent confidence high	SLOT_FILLING	Extract entities, check requirements
SLOT_FILLING	—	all required slots filled	TOOL_CALL	Call relevant tool/API
SLOT_FILLING	—	slot(s) missing	CLARIFICATION	Ask for clarification, follow-up questions
CLARIFICATION	user reply	slot filled	SLOT_FILLING	Re-check requirements, complete
TOOL_CALL	—	tool executes	RESPONSE_GENERATION	Format output, generate natural response
RESPONSE_GENERATION	—	—	INTENT_DETECTION	Wait for user to ask another question (multi-loop)
INTENT_DETECTION	any utterance	out-of-domain	OUT_OF_SCOPE	Politely redirect to support topics
INTENT_DETECTION	any utterance	sensitive/unsafe keywords	SAFETY_ESCALATION	Safety response

Current State	User Input Type	Condition	Next State	Sys Act
				escalat humar
INTENT_DETECTION	goodbye/thanks	—	END	Close conver politel

1.8 Possible Failure Cases

- **Ambiguity:** "Tell me about the exam" — which exam, which student? Requires clarification.
- **Missing slot values:** User asks for fee status without providing student ID.
- **Irrelevant / out-of-domain queries:** Weather, sports scores, general chit-chat unrelated to the university.
- **Unsafe requests:** Self-harm/mental-health distress signals, requests for other students' private data, abusive language, or requests to bypass academic integrity rules (e.g., "give me exam answers").
- **Contradictory input:** User first says "I'm a UG student" then later says "I'm doing my PhD" in the same session.

1.9 Explanation — Why Dialogue Systems Need Intent Detection, Entity Extraction, Context Tracking, and State Management

A task-oriented conversational AI must convert unstructured natural language into structured, actionable information:

- **Intent detection** identifies *what* the user wants to accomplish (e.g., check a fee vs. register a complaint) — without it, the system cannot decide which downstream logic/tool to invoke.
- **Entity extraction** identifies the *specific parameters* needed to fulfil that intent (e.g., which student, which exam, which date) — intents alone are not actionable without their arguments.
- **Context/state tracking** is essential because real conversations span multiple turns; users rarely provide all information in one sentence. Without state tracking, the bot would ask redundant questions, lose track of previously given information, and feel robotic and frustrating.
- **State management** (formal dialogue states like `SLOT_FILLING`, `CLARIFICATION`, `TOOL_CALL`) gives the system a principled way to decide the *next best action* at every turn, rather than reacting to each utterance in isolation.

Why University Helpdesk benefits from conversational automation: Helpdesk queries are high-volume, highly repetitive (fee dates, exam schedules, hostel status), and time-sensitive, especially around admission/exam seasons. Automating them frees human staff for complex, empathetic, or escalated cases while giving students instant 24×7 access to information.

1.10 Inference

Intents such as `fee_inquiry`, `exam_schedule_inquiry`, `library_inquiry`, and `course_inquiry` are **highly automatable** because they map to deterministic lookups in structured backend data with low ambiguity and no emotional/safety sensitivity. In contrast, `complaint_registration` (especially complaints about faculty behaviour) and any conversation containing **safety-relevant signals** (distress, harassment, abuse) are **poor candidates for full automation** — they require empathetic handling, judgment, and human escalation, and the bot's role should be limited to safely logging the issue and routing it to a human, not resolving it autonomously.

Task 2: Intent Detection, Entity Extraction, and Dialogue State Tracking (2 Marks)

In this task we build:

1. A labelled corpus of 80+ utterances covering 12 intents.
2. A **TF-IDF + Logistic Regression** intent classifier (chosen for interpretability, low data requirements, and strong performance on short-text classification — an appropriate middle ground between brittle rule-based matching and heavyweight transformer models that need large fine-tuning data and GPU resources not guaranteed in the Virtual Lab).
3. A **regex + lookup-table based entity extractor** (chosen because helpdesk entities like IDs, dates, emails, and course/department names follow fairly regular, enumerable patterns — regex/gazetteer extraction is fast, fully interpretable, and does not require training data or a GPU, unlike a full NER model).
4. A **Dialogue State Tracker (DST)** class that accumulates slots across turns, tracks what's filled vs. missing, and updates dynamically.

```
In [2]: # =====
# TASK 2.1: Labelled utterance corpus (>= 50 utterances, 12 intents)
# =====
labelled_utterances = [
    # greeting
    ("Hi there", "greeting"),
    ("Hello, good morning", "greeting"),
    ("Hey, is anyone available to help?", "greeting"),
    ("Good afternoon, I need some help", "greeting"),
    ("Hi bot, are you there?", "greeting"),
    ("Hello", "greeting"),
    # course_inquiry
    ("Can you tell me about the B.Tech Computer Science course?", "course_inquiry"),
    ("What is the eligibility for MBA admission?", "course_inquiry"),
    ("How many credits does the Data Science program have?", "course_inquiry"),
    ("Tell me the duration of the M.Tech course", "course_inquiry"),
    ("What subjects are covered in the Computer Science department?", "course_inquiry"),
    ("Is there a minor in AI available for engineering students?", "course_inquiry"),
    ("What is the course structure for B.Sc Mathematics?", "course_inquiry"),
    # fee_inquiry
```

```

("What is my hostel fee for this semester?", "fee_inquiry"),
("How much is the tuition fee for semester 3?", "fee_inquiry"),
("When is the exam fee due?", "fee_inquiry"),
("Can you tell me the total fee for my course?", "fee_inquiry"),
("I want to know the fee structure for B.Tech", "fee_inquiry"),
("Is there a late fee if I pay after the due date?", "fee_inquiry"),
("What's the hostel fee amount for student 21BCE1042?", "fee_inquiry"),
# exam_schedule_inquiry
("When is my DBMS exam?", "exam_schedule_inquiry"),
("What is the schedule for the mid-semester exams?", "exam_schedule_inquiry"),
("Can you tell me the end-semester exam dates?", "exam_schedule_inquiry"),
("Is the Operating Systems exam on 20th September?", "exam_schedule_inquiry"),
("Where will my exam be held?", "exam_schedule_inquiry"),
("What time does the Maths exam start?", "exam_schedule_inquiry"),
("I need the exam timetable for semester 4", "exam_schedule_inquiry"),
# admission_status_inquiry
("What is the status of my admission application?", "admission_status_inquir
("Has my application been approved yet?", "admission_status_inquiry"),
("I applied for MBA, can you check my admission status?", "admission_status_
("Did I get selected for the Computer Science program?", "admission_status_i
("Can you check if my documents were verified for admission?", "admission_st
("My application ID is APP2026123, what's the status?", "admission_status_in
# hostel_inquiry
("Is a room available in Block A hostel?", "hostel_inquiry"),
("I want to know my hostel room allotment", "hostel_inquiry"),
("Can I change my hostel block?", "hostel_inquiry"),
("What is the hostel occupancy for Block C?", "hostel_inquiry"),
("How do I apply for hostel accommodation?", "hostel_inquiry"),
("Are single rooms available in the girls hostel?", "hostel_inquiry"),
# library_inquiry
("Is the book Introduction to Algorithms available in the library?", "librar
("When is my library book due?", "library_inquiry"),
("Can I renew my library book online?", "library_inquiry"),
("How many books can I borrow at once?", "library_inquiry"),
("Do you have Clean Code by Robert Martin in the library?", "library_inquiry
("I want to check the library fine on my account", "library_inquiry"),
# complaint_registration
("I want to file a complaint about hostel maintenance", "complaint_registrat
("There is no water supply in Block B hostel, please register a complaint",
("I want to complain about a faculty member's behaviour", "complaint_registr
("My complaint about mess food quality was not resolved", "complaint_registr
("Please register a grievance regarding the internet connectivity in hostel"
("I need to raise a ticket for a broken washroom in the hostel", "complaint_
# scholarship_inquiry
("Am I eligible for the merit scholarship?", "scholarship_inquiry"),
("What is the status of my scholarship application?", "scholarship_inquiry")
("How much amount does the sports scholarship cover?", "scholarship_inquiry")
("Can I apply for a need-based scholarship this semester?", "scholarship_inq
("What documents are required for the scholarship application?", "scholarshi
# faculty_contact_inquiry
("Can you give me Dr. Mehta's email address?", "faculty_contact_inquiry"),
("What are the office hours of Professor Sharma?", "faculty_contact_inquiry")
("I need the contact number of the Computer Science HOD", "faculty_contact_i
("How do I contact my academic advisor?", "faculty_contact_inquiry"),
("Who is the faculty in charge of the AI course?", "faculty_contact_inquiry")
# thanks_goodbye
("Thank you so much for your help", "thanks_goodbye"),
("That's all I needed, bye!", "thanks_goodbye"),
("Thanks, that was helpful", "thanks_goodbye"),
("Okay great, goodbye", "thanks_goodbye"),

```

```

    ("Perfect, thank you and see you", "thanks_goodbye"),
    # out_of_scope
    ("What's the weather like tomorrow?", "out_of_scope"),
    ("Can you tell me the latest cricket score?", "out_of_scope"),
    ("Recommend me a good movie to watch", "out_of_scope"),
    ("What is the capital of France?", "out_of_scope"),
    ("Can you help me book a flight ticket?", "out_of_scope"),
    ("Tell me a joke", "out_of_scope"),
]

df_utterances = pd.DataFrame(labelled_utterances, columns=["utterance", "intent"])
print(f"Total labelled utterances: {len(df_utterances)}")
print(f"Number of intents: {df_utterances['intent'].nunique()}")
display(df_utterances["intent"].value_counts())

```

Total labelled utterances: 72

Number of intents: 12

```

intent
course_inquiry      7
fee_inquiry         7
exam_schedule_inquiry 7
greeting            6
admission_status_inquiry 6
hostel_inquiry      6
library_inquiry     6
complaint_registration 6
out_of_scope        6
scholarship_inquiry 5
faculty_contact_inquiry 5
thanks_goodbye      5
Name: count, dtype: int64

```

```

In [3]: # =====
# TASK 2.2: Intent Classification -- TF-IDF + Logistic Regression
# =====
X = df_utterances["utterance"]
y = df_utterances["intent"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42, stratify=y
)

vectorizer = TfidfVectorizer(ngram_range=(1, 2), lowercase=True, stop_words="eng")
X_train_vec = vectorizer.fit_transform(X_train)
X_test_vec = vectorizer.transform(X_test)

intent_clf = LogisticRegression(max_iter=1000, C=5.0)
intent_clf.fit(X_train_vec, y_train)

y_pred = intent_clf.predict(X_test_vec)
print("Held-out test accuracy:", round(accuracy_score(y_test, y_pred), 3))
print("\nClassification Report (held-out split):\n")
print(classification_report(y_test, y_pred, zero_division=0))

def predict_intent(utterance: str, top_k: int = 1):
    """Returns predicted intent and confidence score for a single utterance."""
    vec = vectorizer.transform([utterance])
    proba = intent_clf.predict_proba(vec)[0]
    classes = intent_clf.classes_

```

```

ranked = sorted(zip(classes, proba), key=lambda x: -x[1])
return ranked[:top_k]

# Quick sanity check
for test_sent in ["What is my hostel fee?", "When is my DBMS exam?", "Tell me a
pred = predict_intent(test_sent)
print(f"{test_sent!r:45} -> {pred[0][0]:28} (confidence={pred[0][1]:.2f})")

```

Held-out test accuracy: 0.778

Classification Report (held-out split):

	precision	recall	f1-score	support
admission_status_inquiry	0.50	1.00	0.67	1
complaint_registration	1.00	1.00	1.00	1
course_inquiry	0.50	0.50	0.50	2
exam_schedule_inquiry	1.00	1.00	1.00	2
faculty_contact_inquiry	0.00	0.00	0.00	1
fee_inquiry	1.00	1.00	1.00	2
greeting	1.00	0.50	0.67	2
hostel_inquiry	1.00	1.00	1.00	2
library_inquiry	1.00	1.00	1.00	2
out_of_scope	0.50	1.00	0.67	1
scholarship_inquiry	1.00	1.00	1.00	1
thanks_goodbye	0.00	0.00	0.00	1
accuracy			0.78	18
macro avg	0.71	0.75	0.71	18
weighted avg	0.78	0.78	0.76	18

```

'What is my hostel fee?'          -> fee_inquiry          (co
nfidence=0.51)
'When is my DBMS exam?'          -> exam_schedule_inquiry      (co
nfidence=0.64)
'Tell me a joke'                 -> out_of_scope          (co
nfidence=0.57)

```

```

In [4]: # =====
# TASK 2.3: Entity Extraction (regex + gazetteer/lookup based)
# =====

KNOWN_COURSES = ["b.tech computer science", "b.tech", "m.tech", "mba", "b.sc mat
"data science", "computer science", "artificial intelligence",
KNOWN_DEPARTMENTS = ["computer science", "mathematics", "electronics", "mechanic
"civil", "electrical", "management"]
KNOWN_FEE_TYPES = ["hostel fee", "tuition fee", "exam fee", "library fee", "late
KNOWN_COMPLAINT_TYPES = ["hostel maintenance", "faculty behaviour", "mess food q
"internet connectivity", "washroom", "water supply", "r
KNOWN_SCHOLARSHIPS = ["merit scholarship", "sports scholarship", "need-based sch
"minority scholarship"]
KNOWN_EXAMS = ["dbms", "operating systems", "maths", "mathematics", "data struct
"mid-semester", "end-semester", "end-sem", "mid-sem"]
KNOWN_BOOKS = ["introduction to algorithms", "clean code", "data structures and
"operating system concepts"]

ENTITY_PATTERNS = {
    "student_id": r"\b[0-9]{2}[A-Za-z]{2,4}[0-9]{3,4}\b",
    "email": r"\b[\w.+-]+@[ \w-]+\.[ \w.-]+\b",

```

```

"phone_number": r"\b[6-9]\d{9}\b",
"date": r"\b\d{1,2}[/-]\d{1,2}[/-]\d{2,4}\b|\btomorrow\b|\btoday\b|\bnext we",
"application_id": r"\bAPP\d{4,10}\b",
"hostel_block": r"\bblock\s+[A-Za-z]\b",
"semester": r"\bsemester\s*\d\b|\b\d(?:st|nd|rd|th)\s+semester\b|\bsem\s*\d\b",
"faculty_name": r"\b(?:Dr\.|Prof\.|Professor)\s+[A-Z][a-zA-Z]+\b",
}

def _lookup_entities(text_lower, vocab_list, label):
    found = [term for term in vocab_list if term in text_lower]
    return found

def extract_entities(utterance: str) -> dict:
    """Rule-based entity extractor combining regex patterns and domain gazetteer"""
    entities = {}
    text_lower = utterance.lower()

    for label, pattern in ENTITY_PATTERNS.items():
        matches = re.findall(pattern, utterance, flags=re.IGNORECASE)
        if matches:
            entities[label] = matches[0] if len(matches) == 1 else matches

    course_hits = _lookup_entities(text_lower, KNOWN_COURSES, "course_name")
    if course_hits:
        entities["course_name"] = max(course_hits, key=len)

    dept_hits = _lookup_entities(text_lower, KNOWN_DEPARTMENTS, "department")
    if dept_hits:
        entities["department"] = dept_hits[0]

    fee_hits = _lookup_entities(text_lower, KNOWN_FEE_TYPES, "fee_type")
    if fee_hits:
        entities["fee_type"] = fee_hits[0]

    complaint_hits = _lookup_entities(text_lower, KNOWN_COMPLAINT_TYPES, "complaint_type")
    if complaint_hits:
        entities["complaint_type"] = complaint_hits[0]

    scholarship_hits = _lookup_entities(text_lower, KNOWN_SCHOLARSHIPS, "scholarship_name")
    if scholarship_hits:
        entities["scholarship_name"] = scholarship_hits[0]

    exam_hits = _lookup_entities(text_lower, KNOWN_EXAMS, "exam_name")
    if exam_hits:
        entities["exam_name"] = exam_hits[0]

    book_hits = _lookup_entities(text_lower, KNOWN_BOOKS, "book_title")
    if book_hits:
        entities["book_title"] = book_hits[0]

    return entities

# Quick sanity checks
sample_tests = [
    "It's 21BCE1042, hostel fee for block A please",
    "My email is ravi.kumar@univ.edu and phone is 9876543210",
    "I want to know about DBMS exam on 20/09/2026",

```

```

    "My application ID is APP2026123",
]
for s in sample_tests:
    print(f"{s!r}\n -> {extract_entities(s)}\n")

```

"It's 21BCE1042, hostel fee for block A please"

```
-> {'student_id': '21BCE1042', 'hostel_block': 'block A', 'fee_type': 'hostel fee'}
```

'My email is ravi.kumar@univ.edu and phone is 9876543210'

```
-> {'email': 'ravi.kumar@univ.edu', 'phone_number': '9876543210'}
```

'I want to know about DBMS exam on 20/09/2026'

```
-> {'date': '20/09/2026', 'exam_name': 'dbms'}
```

'My application ID is APP2026123'

```
-> {'application_id': 'APP2026123'}
```

```

In [5]: # =====
# TASK 2.4: Dialogue State Tracking (DST)
# =====

REQUIRED_SLOTS = {
    "greeting": [],
    "course_inquiry": ["course_name"],
    "fee_inquiry": ["student_id", "fee_type"],
    "exam_schedule_inquiry": ["student_id", "exam_name"],
    "admission_status_inquiry": ["application_id"],
    "hostel_inquiry": ["student_id"],
    "library_inquiry": ["book_title"],
    "complaint_registration": ["student_id", "complaint_type"],
    "scholarship_inquiry": ["student_id"],
    "faculty_contact_inquiry": ["faculty_name"],
    "thanks_goodbye": [],
    "out_of_scope": [],
}

class DialogueStateTracker:
    """Maintains dialogue state (intent + slots) across multiple conversation turns"""

    def __init__(self, session_id):
        self.session_id = session_id
        self.current_intent = None
        self.slots = {}
        self.turn_log = []

    def get_missing_slots(self):
        if self.current_intent is None:
            return []
        required = REQUIRED_SLOTS.get(self.current_intent, [])
        return [s for s in required if s not in self.slots or not self.slots[s]]

    def update(self, utterance: str):
        entities = extract_entities(utterance)
        pred_intent, conf = predict_intent(utterance)[0]
        missing_before = self.get_missing_slots()
        slot_fill_match = any(k in missing_before for k in entities.keys())

```

```

        if self.current_intent is None:
            self.current_intent = pred_intent
        elif not missing_before:
            self.current_intent = pred_intent
            self.slots = {}
        elif slot_fill_match or conf < 0.5:
            pass # continuation turn: keep current task, just add new slot value
        else:
            self.current_intent = pred_intent
            self.slots = {}

        self.slots.update(entities)
        missing_after = self.get_missing_slots()

        if self.current_intent in ("greeting", "thanks_goodbye"):
            next_action = "Respond politely / " + ("greet user" if self.current_intent == "greeting" else "thank user")
        elif self.current_intent == "out_of_scope":
            next_action = "Politely redirect to supported topics"
        elif missing_after:
            next_action = f"Ask clarifying question for: {' '.join(missing_after)}"
        else:
            next_action = f"Call tool for intent '{self.current_intent}' and generate response"

        record = {
            "utterance": utterance,
            "predicted_intent": pred_intent,
            "confidence": round(conf, 2),
            "extracted_entities": entities,
            "dialogue_state": dict(self.slots),
            "missing_information": missing_after if missing_after else "None",
            "next_system_action": next_action,
        }
        self.turn_log.append(record)
        return record

print("DialogueStateTracker class defined successfully.")

```

DialogueStateTracker class defined successfully.

```

In [6]: # =====
# TASK 2.5: Test DST on 10 multi-turn conversations
# =====
test_conversations = {
    "Conv-1 (Fee Inquiry)": [
        "Hi, I want to know my hostel fee for this semester",
        "It's 21BCE1042",
    ],
    "Conv-2 (Exam Schedule)": [
        "When is my exam?",
        "DBMS, my ID is 22BCS0091",
    ],
    "Conv-3 (Hostel Inquiry)": [
        "Is a room available in block A?",
        "My student ID is 20BME0033",
    ],
    "Conv-4 (Library Inquiry)": [
        "Is the book Introduction to Algorithms available in the library?",
    ],
    "Conv-5 (Complaint Registration)": [

```

```

        "I want to file a complaint about hostel maintenance",
        "My student ID is 21BCE1042",
    ],
    "Conv-6 (Scholarship Inquiry)": [
        "Am I eligible for the merit scholarship?",
        "My student ID is 23BCS0087",
    ],
    "Conv-7 (Admission Status)": [
        "What is the status of my admission application?",
        "My application ID is APP2026123",
    ],
    "Conv-8 (Course Inquiry)": [
        "Can you tell me about the B.Tech Computer Science course?",
    ],
    "Conv-9 (Faculty Contact)": [
        "Can you give me Dr. Mehta's email address?",
    ],
    "Conv-10 (Greeting + Fee + Goodbye)": [
        "Hello",
        "What is the tuition fee for semester 3, my ID is 22BCS0091",
        "Thank you, that's all",
    ],
}

task2_rows = []
for conv_name, turns in test_conversations.items():
    dst = DialogueStateTracker(conv_name)
    for turn in turns:
        rec = dst.update(turn)
        task2_rows.append({
            "Conversation": conv_name,
            "Utterance": rec["utterance"],
            "Predicted Intent": f"{rec['predicted_intent']} ({rec['confidence']})",
            "Extracted Entity": rec["extracted_entities"] if rec["extracted_entities"] else "",
            "Current Dialogue State": rec["dialogue_state"] if rec["dialogue_state"] else "",
            "Missing Information": rec["missing_information"] if rec["missing_information"] else "",
            "Next System Action": rec["next_system_action"] if rec["next_system_action"] else "",
        })

df_task2 = pd.DataFrame(task2_rows)
display(df_task2)

```

	Conversation	Utterance	Predicted Intent	Extracted Entity	Current Dialog State
0	Conv-1 (Fee Inquiry)	Hi, I want to know my hostel fee for this semester	fee_inquiry (0.5)	{'fee_type': 'hostel fee'}	{'fee_type': 'hostel fee'}
1	Conv-1 (Fee Inquiry)	It's 21BCE1042	out_of_scope (0.1)	{'student_id': '21BCE1042'}	{'fee_type': 'hostel fee', 'student_id': '21BCE1042'}
2	Conv-2 (Exam Schedule)	When is my exam?	exam_schedule_inquiry (0.62)	-	-
3	Conv-2 (Exam Schedule)	DBMS, my ID is 22BCS0091	exam_schedule_inquiry (0.18)	{'student_id': '22BCS0091', 'exam_name': 'dbms'}	{'student_id': '22BCS0091', 'exam_name': 'dbms'}
4	Conv-3 (Hostel Inquiry)	Is a room available in block A?	hostel_inquiry (0.25)	{'hostel_block': 'block A'}	{'hostel_block': 'block A'}
5	Conv-3 (Hostel Inquiry)	My student ID is 20BME0033	admission_status_inquiry (0.19)	{'student_id': '20BME0033'}	{'hostel_block': 'block A', 'student_id': '20BME0033'}
6	Conv-4 (Library Inquiry)	Is the book Introduction to Algorithms available in the library?	library_inquiry (0.39)	{'book_title': 'introduction to algorithms'}	{'book_title': 'introduction to algorithms'}
7	Conv-5 (Complaint Registration)	I want to file a complaint about hostel maintenance	complaint_registration (0.61)	{'complaint_type': 'hostel maintenance'}	{'complaint_type': 'hostel maintenance'}
8	Conv-5 (Complaint Registration)	My student ID is 21BCE1042	admission_status_inquiry (0.19)	{'student_id': '21BCE1042'}	{'complaint_type': 'hostel maintenance', 'student_id': '21BCE1042'}
9	Conv-6 (Scholarship Inquiry)	Am I eligible for the merit scholarship?	scholarship_inquiry (0.56)	{'scholarship_name': 'merit scholarship'}	{'scholarship_name': 'merit scholarship'}
10	Conv-6 (Scholarship Inquiry)	My student ID is 23BCS0087	admission_status_inquiry (0.19)	{'student_id': '23BCS0087'}	{'scholarship_name': 'merit scholarship', 'student_id': '23BCS0087'}
11	Conv-7 (Admission Status)	What is the status of my admission application?	admission_status_inquiry (0.66)	-	-

	Conversation	Utterance	Predicted Intent	Extracted Entity	Current Dialog State
12	Conv-7 (Admission Status)	My application ID is APP2026123	admission_status_inquiry (0.44)	{'application_id': 'APP2026123'}	{'application_id': 'APP2026123'}
13	Conv-8 (Course Inquiry)	Can you tell me about the B.Tech Computer Science course?	course_inquiry (0.61)	{'course_name': 'b.tech computer science', 'department': 'computer science'}	{'course_name': 'b.tech computer science', 'department': 'computer science'}
14	Conv-9 (Faculty Contact)	Can you give me Dr. Mehta's email address?	faculty_contact_inquiry (0.54)	{'faculty_name': 'Dr. Mehta'}	{'faculty_name': 'Dr. Mehta'}
15	Conv-10 (Greeting + Fee + Goodbye)	Hello	greeting (0.2)	-	-
16	Conv-10 (Greeting + Fee + Goodbye)	What is the tuition fee for semester 3, my ID is 22BCS0091	fee_inquiry (0.6)	{'student_id': '22BCS0091', 'semester': 'semester 3', 'fee_type': 'tuition fee'}	{'student_id': '22BCS0091', 'semester': 'semester 3', 'fee_type': 'tuition fee'}
17	Conv-10 (Greeting + Fee + Goodbye)	Thank you, that's all	thanks_goodbye (0.5)	-	-

Task 2 — Explanation, Justification, and Inference

Logic used: The utterance first passes through a TF-IDF vectorizer (unigrams + bigrams) and a Logistic Regression classifier to predict the intent with a confidence score. In parallel, a rule-based extractor (regex for structured patterns like IDs/emails/dates + gazetteer lookups for domain vocabulary like course/department/fee-type names) pulls out entities. The `DialogueStateTracker` then merges newly extracted entities into a running `slots` dictionary tied to the **current active intent**, and only resets the state once the active task's required slots are all filled (task completed) or a new intent is detected with high confidence.

Justification: TF-IDF + Logistic Regression is justified over a purely rule-based intent matcher because it generalizes to paraphrases not seen verbatim in training (e.g., "how much is my tuition" vs. "what is my tuition fee"), while remaining fast and fully explainable (feature weights are inspectable) — appropriate given our modestly sized, domain-specific corpus. A transformer/LLM approach was considered but rejected for the core classifier because (a) it needs far more labelled data or heavy prompting

infrastructure than is necessary for a bounded 12-intent problem, and (b) interpretability and low-latency, offline execution are more important in a helpdesk setting than marginal accuracy gains.

How DST avoids repeating questions: Because slots are persisted across turns for the *same* active intent, when a user supplies only their student ID in turn 2 (e.g., Conv-2: "DBMS, my ID is 22BCS0091"), the tracker does not re-ask "what do you want help with?" — it recognizes the utterance as a slot-fill for the still-open `exam_schedule_inquiry` task, merges the `student_id` and `exam_name` entities into the same state object, and immediately determines that all required slots are now filled, moving to `Call tool...`. This is the essence of coherent multi-turn interaction: the system's next action is always a function of the **accumulated state**, not just the latest isolated utterance.

Inference: Across the 10 test conversations, the classifier correctly identified the primary intent whenever the first turn of a conversation contained clear task-indicating vocabulary (e.g., "fee", "exam", "scholarship", "complaint"). Follow-up turns that consisted purely of entity values (e.g., "It's 21BCE1042") were sometimes mis-classified in isolation by the intent model, but the DST's slot-fill heuristic correctly retained the original active intent rather than incorrectly switching tasks — demonstrating that **state management compensates for the intent classifier's weaknesses on short, entity-only utterances**, and confirming that context tracking is not optional but a structural necessity for multi-turn task completion.

Task 3: Tool-Augmented Response Generation and Action Execution (2 Marks)

We implement **8 simulated backend tools/APIs** (exceeding the minimum of 3) backed by in-memory "databases", and a **tool-routing layer** that decides — for every user query — whether a tool is needed, which tool to call, what parameters it needs, and how to turn the raw tool output into a natural-language response.

```
In [7]: # =====
# TASK 3.1: Simulated backend databases
# =====
STUDENTS_DB = {
    "21BCE1042": {"name": "Ravi Kumar", "semester": 3, "department": "Computer S
                "fees": {"hostel fee": {"amount": 45000, "due_date": "30/08/20
                        "tuition fee": {"amount": 120000, "due_date": "15/08/
                        "exam fee": {"amount": 2000, "due_date": "01/09/2026"
    "22BCS0091": {"name": "Sneha Iyer", "semester": 4, "department": "Computer S
                "fees": {"hostel fee": {"amount": 47500, "due_date": "01/09/20
                        "tuition fee": {"amount": 125000, "due_date": "20/08/
    "20BME0033": {"name": "Arjun Nair", "semester": 5, "department": "Mechanical
                "fees": {"hostel fee": {"amount": 46000, "due_date": "28/08/20
    "23BCS0087": {"name": "Priya Das", "semester": 1, "department": "Computer Sc
                "fees": {"tuition fee": {"amount": 118000, "due_date": "10/08/
}

EXAM_SCHEDULE_DB = {
```

```

    "dbms": {"date": "20/09/2026", "time": "10:00 AM", "hall": "Hall 3"},
    "operating systems": {"date": "22/09/2026", "time": "02:00 PM", "hall": "Hall 3"},
    "data structures": {"date": "24/09/2026", "time": "10:00 AM", "hall": "Hall 3"},
    "maths": {"date": "26/09/2026", "time": "09:00 AM", "hall": "Hall 4"},
    "mid-semester": {"date": "05/09/2026", "time": "09:00 AM", "hall": "Various"}
}

HOSTEL_DB = {
    "block a": {"total_rooms": 100, "occupied": 92},
    "block b": {"total_rooms": 80, "occupied": 80},
    "block c": {"total_rooms": 60, "occupied": 40},
}

LIBRARY_DB = {
    "introduction to algorithms": {"available": True, "copies": 2},
    "clean code": {"available": False, "due_date": "18/08/2026"},
    "data structures and algorithms": {"available": True, "copies": 1},
    "operating system concepts": {"available": False, "due_date": "22/08/2026"},
}

SCHOLARSHIP_DB = {
    "21BCE1042": {"eligible": True, "scholarship": "Merit Scholarship", "amount": 1000},
    "23BCS0087": {"eligible": True, "scholarship": "Need-Based Scholarship", "amount": 500},
    "20BME0033": {"eligible": False, "reason": "CGPA below cutoff"},
}

ADMISSION_DB = {
    "APP2026123": {"status": "Approved", "course": "MBA"},
    "APP2026456": {"status": "Under Review", "course": "B.Tech Computer Science"},
}

FACULTY_DB = {
    "dr. mehta": {"email": "mehta@univ.edu", "office_hours": "Mon-Wed 2-4 PM"},
    "prof. sharma": {"email": "sharma@univ.edu", "office_hours": "Tue-Thu 11 AM-2 PM"},
}

COMPLAINT_LOG = [] # simulated ticketing system, grows at runtime

# =====
# TASK 3.2: Tool / simulated-API implementations (8 tools)
# =====

def tool_fee_calculator(student_id=None, fee_type=None):
    if not student_id or student_id not in STUDENTS_DB:
        return {"error": "Student ID not found in records."}
    student = STUDENTS_DB[student_id]
    if fee_type and fee_type in student["fees"]:
        info = student["fees"][fee_type]
        return {"student_id": student_id, "fee_type": fee_type, "amount": info["amount"]}
    return {"student_id": student_id, "all_fees": student["fees"]}

def tool_exam_schedule_checker(exam_name=None, student_id=None):
    if not exam_name:
        return {"error": "Exam name not specified."}
    key = exam_name.lower()
    for k, v in EXAM_SCHEDULE_DB.items():
        if k in key or key in k:
            return {"exam_name": exam_name, "date": v["date"], "time": v["time"]}
    return {"error": f"No schedule found for exam '{exam_name}'."}

```

```

def tool_hostel_availability_checker(hostel_block=None, student_id=None):
    if hostel_block:
        block_key = hostel_block.lower().replace("block ", "block ").strip()
        info = HOSTEL_DB.get(block_key)
        if not info:
            return {"error": f"No data for {hostel_block}."}
        available = info["total_rooms"] - info["occupied"]
        return {"hostel_block": hostel_block, "available_rooms": available, "tot
    if student_id and student_id in STUDENTS_DB:
        return {"student_id": student_id, "hostel_status": "Allotted", "block":
    return {"error": "Insufficient information to check hostel status."}

def tool_library_checker(book_title=None):
    if not book_title:
        return {"error": "Book title not specified."}
    info = LIBRARY_DB.get(book_title.lower())
    if not info:
        return {"error": f"'{book_title}' not found in catalogue."}
    return {"book_title": book_title, **info}

def tool_complaint_registration(student_id=None, complaint_type=None):
    if not complaint_type:
        return {"error": "Complaint type not specified."}
    ticket_id = f"TCK{1000 + len(COMPLAINT_LOG) + 1}"
    ticket = {"ticket_id": ticket_id, "student_id": student_id or "N/A",
              "complaint_type": complaint_type, "status": "Registered"}
    COMPLAINT_LOG.append(ticket)
    return ticket

def tool_scholarship_checker(student_id=None):
    if not student_id or student_id not in SCHOLARSHIP_DB:
        return {"error": "No scholarship record found for this student ID."}
    return {"student_id": student_id, **SCHOLARSHIP_DB[student_id]}

def tool_admission_status_checker(application_id=None):
    if not application_id or application_id not in ADMISSION_DB:
        return {"error": "Application ID not found."}
    return {"application_id": application_id, **ADMISSION_DB[application_id]}

def tool_faculty_lookup(faculty_name=None):
    if not faculty_name or faculty_name.lower() not in FACULTY_DB:
        return {"error": "Faculty not found in directory."}
    return {"faculty_name": faculty_name, **FACULTY_DB[faculty_name.lower()]}

print("8 tools implemented: FeeCalculator, ExamScheduleChecker, HostelAvailabili
print("LibraryChecker, ComplaintRegistrationTool, ScholarshipEligibilityChecker,
print("AdmissionStatusChecker, FacultyContactLookup")

```

8 tools implemented: FeeCalculator, ExamScheduleChecker, HostelAvailabilityChecke
r,
LibraryChecker, ComplaintRegistrationTool, ScholarshipEligibilityChecker,
AdmissionStatusChecker, FacultyContactLookup

```

In [8]: # =====
# TASK 3.3: Tool-selection / routing logic + response generation
# =====
INTENT_TOOL_MAP = {
    "fee_inquiry": ("Fee Calculator", tool_fee_calculator, ["student_id", "fee_t
    "exam_schedule_inquiry": ("Exam Schedule Checker", tool_exam_schedule_checke
    "hostel_inquiry": ("Hostel Availability Checker", tool_hostel_availability_c
    "library_inquiry": ("Library Search Tool", tool_library_checker, ["book_titl
    "complaint_registration": ("Complaint Registration Tool", tool_complaint_reg
    "scholarship_inquiry": ("Scholarship Eligibility Checker", tool_scholarship_
    "admission_status_inquiry": ("Admission Status Checker", tool_admission_stat
    "faculty_contact_inquiry": ("Faculty Contact Lookup", tool_faculty_lookup, [
}

def generate_nl_response(intent, tool_output):
    """Converts structured tool output (dict) into a user-friendly natural-langu
    if "error" in tool_output:
        return f"I'm sorry, {tool_output['error']} Could you please double-check

    if intent == "fee_inquiry":
        if "fee_type" in tool_output:
            return (f"For student {tool_output['student_id']}, the {tool_output[
                f"₹{tool_output['amount']:,}, due on {tool_output['due_date']
            return f"Here are all the fee details for student {tool_output['student_
    if intent == "exam_schedule_inquiry":
        return (f"The {tool_output['exam_name']} exam is scheduled on {tool_outp
            f"at {tool_output['time']} in {tool_output['hall']}."))
    if intent == "hostel_inquiry":
        if "available_rooms" in tool_output:
            return f"{tool_output['hostel_block']} currently has {tool_output['a
            return f"Student {tool_output['student_id']} is allotted {tool_output['b
    if intent == "library_inquiry":
        if tool_output.get("available"):
            return f"'{tool_output['book_title']}' is available in the library (
            return f"'{tool_output['book_title']}' is currently checked out and due
    if intent == "complaint_registration":
        return f"Your complaint about '{tool_output['complaint_type']}' has been
    if intent == "scholarship_inquiry":
        if tool_output.get("eligible"):
            return f"Good news! Student {tool_output['student_id']} is eligible
            return f"Student {tool_output['student_id']} is currently not eligible f
    if intent == "admission_status_inquiry":
        return f"Application {tool_output['application_id']} for {tool_output['c
    if intent == "faculty_contact_inquiry":
        return (f"{tool_output['faculty_name']} ({tool_output['dept']} dept.) ca
            f"{tool_output['email']}. Office hours: {tool_output['office_hou
    return "Here is the information you requested."

def decide_and_call_tool(user_query: str, extra_slots: dict = None):
    """Full tool-augmented pipeline: intent -> tool decision -> parameter extrac
    intent, conf = predict_intent(user_query)[0]
    entities = extract_entities(user_query)
    if extra_slots:
        entities = {**extra_slots, **entities}

    if intent not in INTENT_TOOL_MAP:
        return {

```

```

        "User Query": user_query, "Predicted Intent": intent, "Required Tool": tool_name,
        "Tool Input": "-", "Tool Output": "-",
        "Generated Response": "This is a general/greeting/out-of-scope message"
    }

    tool_name, tool_fn, params = INTENT_TOOL_MAP[intent]
    tool_input = {p: entities.get(p) for p in params if entities.get(p)}
    tool_output = tool_fn(**tool_input) if tool_input else {"error": "insufficient information"}
    response = generate_nl_response(intent, tool_output)

    return {
        "User Query": user_query, "Predicted Intent": intent, "Required Tool": tool_name,
        "Tool Input": tool_input if tool_input else "None provided",
        "Tool Output": tool_output, "Generated Response": response,
    }

# =====
# Demo: run tool-augmented pipeline on 8 representative queries
# =====
task3_queries = [
    "What is my hostel fee? My student ID is 21BCE1042",
    "When is the DBMS exam? My ID is 22BCS0091",
    "Is a room available in block A?",
    "Is the book Clean Code available in the library?",
    "I want to file a complaint about hostel maintenance, my ID is 21BCE1042",
    "Am I eligible for a scholarship? My student ID is 23BCS0087",
    "What is the status of application APP2026123?",
    "Can you give me Dr. Mehta's contact details?",
]

task3_results = [decide_and_call_tool(q) for q in task3_queries]
df_task3 = pd.DataFrame(task3_results)
display(df_task3)

```

	User Query	Predicted Intent	Required Tool	Tool Input	Tool Output
0	What is my hostel fee? My student ID is 21BCE1042	fee_inquiry	Fee Calculator	{'student_id': '21BCE1042', 'fee_type': 'hostel fee'}	{'student_id': '21BCE1042', 'fee_type': 'hostel fee', 'amount': 45000, 'due_date': '30/08/2026'}
1	When is the DBMS exam? My ID is 22BCS0091	exam_schedule_inquiry	Exam Schedule Checker	{'exam_name': 'dbms', 'student_id': '22BCS0091'}	{'exam_name': 'dbms', 'date': '20/09/2026', 'time': '10:00 AM', 'hall': 'Hall 3'}
2	Is a room available in block A?	hostel_inquiry	Hostel Availability Checker	{'hostel_block': 'block A'}	{'hostel_block': 'block A', 'available_rooms': 8, 'total_rooms': 100}
3	Is the book Clean Code available in the library?	library_inquiry	Library Search Tool	{'book_title': 'clean code'}	{'book_title': 'clean code', 'available': False, 'due_date': '18/08/2026'}
4	I want to file a complaint about hostel maintenance, my ID is 21BCE1042	complaint_registration	Complaint Registration Tool	{'student_id': '21BCE1042', 'complaint_type': 'hostel maintenance'}	{'ticket_id': 'TCK1001', 'student_id': '21BCE1042', 'complaint_type': 'hostel maintenance', 'status': 'Registered'}
5	Am I eligible for a scholarship? My student ID is 23BCS0087	scholarship_inquiry	Scholarship Eligibility Checker	{'student_id': '23BCS0087'}	{'student_id': '23BCS0087', 'eligible': True, 'scholarship': 'Need-Based Scholarship', 'amount': 15000}
6	What is the status of application APP2026123?	admission_status_inquiry	Admission Status Checker	{'application_id': 'APP2026123'}	{'application_id': 'APP2026123', 'status': 'Approved', 'course': 'MBA'}
7	Can you give me Dr. Mehta's contact details?	faculty_contact_inquiry	Faculty Contact Lookup	{'faculty_name': 'Dr. Mehta'}	{'faculty_name': 'Dr. Mehta', 'email': 'mehta@univ.edu', 'office_hours': 'Mon-Wed 2-4 PM', 'dept': 'Computer Science'}

Task 3 — Explanation, Justification, and Inference

Logic used: For every query, the pipeline (1) predicts the intent, (2) checks a static `INTENT_TOOL_MAP` to decide whether a tool is required and which one, (3) extracts only the parameters that specific tool needs from the entity dictionary, (4) calls the tool function (a stand-in for a real database/REST API call), and (5) passes the structured JSON-like output through an intent-specific natural-language template to produce the final response.

Justification: Separating *decision* (which tool), *execution* (calling the tool), and *presentation* (NL templating) mirrors how production tool-augmented LLM/agent systems are built (e.g., function-calling architectures), and keeps each responsibility independently testable. Templated response generation is used instead of free-form generation because helpdesk answers must be **factually exact** (fee amounts, dates, ticket IDs) — a generative model paraphrasing these numbers risks hallucinating incorrect figures, whereas a template guarantees the tool's ground-truth values are reproduced verbatim.

How tool use improves factuality and task completion vs. pure text generation: A pure text-generation chatbot (e.g., an LLM with no tool access) can only produce an answer based on patterns learned during training or invent a plausible-sounding number/date — it has no way to know student 21BCE1042's *actual* current hostel fee or the *actual* real-time availability of a library book. By grounding every factual claim in a deterministic tool call to a system of record, the response is guaranteed to be accurate and up to date, and every claim in the response is *traceable* back to a specific tool call and its output. This also directly enables **task completion** (e.g., actually registering a complaint and returning a real ticket ID), which pure text generation cannot do at all since it has no side effects on the real world.

Inference: In the demo above, all 8 queries containing sufficient entities were correctly routed to their respective tools and produced accurate, verifiable responses (e.g., the ticket ID for the complaint, the exact exam hall/time). This confirms that tool integration **eliminates hallucination risk for factual/transactional queries** and **increases user trust**, because the user can be confident the amount, date, or ticket number quoted is real rather than a plausible guess. The main remaining risk is **incorrect parameter extraction** (e.g., a mis-parsed student ID) silently causing a wrong-but-confident tool response — this motivates the confirmation/clarification safeguards added in Task 4.

Task 4: Memory, Personalization, Ambiguity Handling, and Safety Control (2 Marks)

This task augments the pipeline with:

1. **Short-term memory** — full turn-by-turn history for the current session.

2. **User preference memory** — persists preferences (preferred contact mode, preferred language) *across* sessions for a given user ID.
3. **Ambiguity detection** — flags low-confidence intents and contradictory slot values.
4. **Clarification generation** — targeted follow-up questions.
5. **Safety/escalation rules** — keyword + pattern based detection of self-harm, abusive language, privacy violations, and academic-integrity violations, each mapped to a safe, non-committal response and an escalation flag.

```
In [9]: # =====
# TASK 4.1: Safety / escalation rules
# =====
SAFETY_PATTERNS = {
    "self_harm_or_distress": [
        "suicide", "kill myself", "end my life", "hopeless", "self harm", "self-
    ],
    "abusive_language": [
        "idiot", "stupid bot", "shut up", "fuck", "useless bot", "damn you"
    ],
    "privacy_violation": [
        "phone number of another student", "share other student", "private data
    ],
    "academic_integrity_violation": [
        "exam answers", "leak the question paper", "cheat in exam", "give me the
    ],
}

SAFE_RESPONSES = {
    "self_harm_or_distress": ("I'm really sorry you're feeling this way. I'm not
        "mental-health support, but please reach out to t
        "1800-XXX-XXXX or a trusted person immediately. W
    "abusive_language": ("I understand you might be frustrated, and I want to he
        "request so I can assist you better?"),
    "privacy_violation": ("I'm not able to share another student's personal info
        "privacy policy. If you have a legitimate, authorized
    "academic_integrity_violation": ("I can't help with that request as it goes
        "I'd be happy to help you with the officia
}

def detect_safety_issue(utterance: str):
    text_lower = utterance.lower()
    for category, patterns in SAFETY_PATTERNS.items():
        if any(p in text_lower for p in patterns):
            return category
    return None

# =====
# TASK 4.2: Ambiguity detection + clarification generation
# =====
def detect_ambiguity(utterance: str, intent: str, confidence: float, entities: d
    if confidence < 0.35:
        return "low_confidence_intent"
    if intent not in ("greeting", "thanks_goodbye", "out_of_scope") and not enti
        return "vague_request_no_entities"
    return None
```

```

def generate_clarification(issue_type: str, intent: str = None):
    if issue_type == "low_confidence_intent":
        return "I'm not fully sure what you need help with. Could you tell me if
    if issue_type == "vague_request_no_entities":
        return f"Sure, I can help with that. Could you share a few more details
    if issue_type == "contradictory_input":
        return "I noticed you gave two different values earlier in this conversa
    return "Could you please provide a bit more detail?"

# =====
# TASK 4.3: Memory -- short-term (per-session) + user preference (cross-session)
# =====
USER_PREFERENCE_STORE = {} # persists across sessions, keyed by user_id

PREFERENCE_PATTERNS = {
    "preferred_contact_mode": r"\b(?:contact|reach)\s+me\s+(?:via|by)\s+(email|p
    "preferred_language": r"\breply\s+in\s+(english|hindi|tamil)\b",
}

class ConversationMemory:
    """Combines short-term (session) memory with cross-session user-preference m
    plus ambiguity/contradiction tracking on top of the Task-2 DialogueStateTrac

    def __init__(self, user_id: str):
        self.user_id = user_id
        self.dst = DialogueStateTracker(session_id=user_id)
        self.short_term_history = [] # list of {user, bot} per turn -- current
        self.preferences = USER_PREFERENCE_STORE.setdefault(user_id, {})
        self._slot_history = {} # tracks previously seen values per slot, for c

    def _update_preferences(self, utterance):
        text_lower = utterance.lower()
        for pref_key, pattern in PREFERENCE_PATTERNS.items():
            m = re.search(pattern, text_lower)
            if m:
                self.preferences[pref_key] = m.group(1)

    def _check_contradiction(self, entities):
        for stable_slot in ["student_id", "application_id"]:
            if stable_slot in entities:
                new_val = entities[stable_slot]
                old_val = self._slot_history.get(stable_slot)
                if old_val and old_val != new_val:
                    return stable_slot, old_val, new_val
                self._slot_history[stable_slot] = new_val
        return None, None, None

    def process_turn(self, utterance: str):
        self._update_preferences(utterance)

        safety_issue = detect_safety_issue(utterance)
        if safety_issue:
            response = SAFE_RESPONSES[safety_issue]
            record = {"User Input": utterance, "Detected Issue": f"Unsafe: {safe
                    "System Decision": "Escalate / Safe-response (no tool cal
                    "Clarification / Safe Response": response,
                    "Updated Memory": {"preferences": dict(self.preferences)}}

```

```

        self.short_term_history.append(record)
        return record

    intent, conf = predict_intent(utterance)[0]
    entities = extract_entities(utterance)

    contra_slot, old_val, new_val = self._check_contradiction(entities)
    if contra_slot:
        response = generate_clarification("contradictory_input")
        record = {"User Input": utterance, "Detected Issue": f"Contradictory",
                  "System Decision": "Pause task, request confirmation",
                  "Clarification / Safe Response": response,
                  "Updated Memory": {"preferences": dict(self.preferences)},
                  self.short_term_history.append(record)
        return record

    ambiguity = detect_ambiguity(utterance, intent, conf, entities)
    if ambiguity:
        response = generate_clarification(ambiguity, intent)
        record = {"User Input": utterance, "Detected Issue": ambiguity,
                  "System Decision": "Ask clarifying question",
                  "Clarification / Safe Response": response,
                  "Updated Memory": {"preferences": dict(self.preferences)}}
        self.short_term_history.append(record)
        return record

    dst_record = self.dst.update(utterance)
    if dst_record["missing_information"] != "None":
        response = generate_clarification("vague_request_no_entities", intent,
        f"Could you also provide: {'', ' '.join(dst_record['missing_informa
        decision = "Ask targeted slot-filling question"
    else:
        if intent in INTENT_TOOL_MAP:
            tool_result = decide_and_call_tool(utterance, extra_slots=self.d
            response = tool_result["Generated Response"]
            decision = f"Call tool: {tool_result['Required Tool']}"
        elif intent == "out_of_scope":
            response = ("I'm a university helpdesk assistant, so I can't hel
            "fees, exams, hostel, admissions, library, scholars
            "would you like help with any of these?")
            decision = "Redirect to supported topics (out of domain)"
        elif intent == "thanks_goodbye":
            response = "You're welcome! Have a great day."
            decision = "Close conversation politely"
        else:
            response = "Hello! How can I help you today?"
            decision = "Direct response (greeting, no tool needed)"

    record = {"User Input": utterance, "Detected Issue": "None",
              "System Decision": decision,
              "Clarification / Safe Response": response,
              "Updated Memory": {"slots": dict(self.dst.slots), "preference
    self.short_term_history.append(record)
    return record

```

```
print("ConversationMemory (short-term + preference memory + ambiguity + safety)
```

ConversationMemory (short-term + preference memory + ambiguity + safety) defined successfully.

```

In [10]: # =====
# TASK 4.4: Test on 5 difficult cases
# =====
task4_rows = []

# Case 1: Missing information
mem1 = ConversationMemory(user_id="user_missing_info")
r1 = mem1.process_turn("When is my exam?")
task4_rows.append({"Case": "1. Missing information", **r1})

# Case 2: Contradictory user input
mem2 = ConversationMemory(user_id="user_contradiction")
mem2.process_turn("What is my hostel fee? My student ID is 21BCE1042")
r2 = mem2.process_turn("Sorry, actually my student ID is 99XYZ0000")
task4_rows.append({"Case": "2. Contradictory user input", **r2})

# Case 3: Ambiguous query
mem3 = ConversationMemory(user_id="user_ambiguous")
r3 = mem3.process_turn("Can you help me?")
task4_rows.append({"Case": "3. Ambiguous query", **r3})

# Case 4: Out-of-domain query
mem4 = ConversationMemory(user_id="user_ood")
r4 = mem4.process_turn("What's the weather like tomorrow?")
task4_rows.append({"Case": "4. Out-of-domain query", **r4})

# Case 5: Sensitive / unsafe query
mem5 = ConversationMemory(user_id="user_sensitive")
r5 = mem5.process_turn("I've been feeling really hopeless about my exam stress 1
task4_rows.append({"Case": "5. Sensitive / unsafe query", **r5})

df_task4 = pd.DataFrame(task4_rows)[["Case", "User Input", "Detected Issue", "Sy
                                     "Clarification / Safe Response", "Updated
display(df_task4)

```

	Case	User Input	Detected Issue	System Decision	Clarification / Safe Response	Update Memory
0	1. Missing information	When is my exam?	vague_request_no_entities	Ask clarifying question	Sure, I can help with that. Could you share a few more details (e.g., your student ID or the specific item you're asking about)?	{'preferen
1	2. Contradictory user input	Sorry, actually my student ID is 99XYZ0000	Contradictory 'student_id': '21BCE1042' vs '99XYZ0000'	Pause task, request confirmation	I noticed you gave two different values earlier in this conversation. Could you confirm which one is correct?	{'preferen 'slot_hist {'student '21BCE10
2	3. Ambiguous query	Can you help me?	low_confidence_intent	Ask clarifying question	I'm not fully sure what you need help with. Could you tell me if it's about fees, exams, hostel, admissions, library, or a complaint?	{'preferen
3	4. Out-of-domain query	What's the weather like tomorrow?	low_confidence_intent	Ask clarifying question	I'm not fully sure what you need help with. Could you tell me if it's about fees, exams, hostel, admissions, library, or a complaint?	{'preferen
4	5. Sensitive / unsafe query	I've been feeling really hopeless about my exam stress lately, I	Unsafe: self_harm_or_distress	Escalate / Safe-response (no tool call)	I'm really sorry you're feeling this way. I'm not equipped to provide medical or mental-	{'preferen

Case	User Input	Detected Issue	System Decision	Clarification / Safe Response	Update Memory
	want to end my life			health support, but please reach out to the campus counselling helpline at 1800-XXX-XXXX or a trusted person immediately. Would you like me to connect you to a human counsellor?	

Task 4 — Explanation, Justification, and Inference

Useful personalization vs. unsafe over-personalization: Useful personalization stores *non-sensitive, task-relevant* preferences the user explicitly volunteers (e.g., "reply in Hindi", "contact me via email") to reduce friction in future interactions — it is opt-in, narrow in scope, and improves usability. **Unsafe over-personalization** would involve silently inferring or storing sensitive attributes (e.g., health status, financial hardship, disciplinary history, emotional state) without consent, retaining them indefinitely, or using them to make consequential decisions (e.g., silently flagging a "distressed" student to faculty without their knowledge). Our design intentionally limits stored preferences to `preferred_contact_mode` and `preferred_language` and never persists anything captured under a safety-triggered turn — safety-relevant disclosures (e.g., self-harm signals) are handled by an immediate safe response, not stored as a "preference".

Why ambiguity handling is essential: Real users under-specify requests constantly ("Can you help me?", "Tell me about the exam") because they assume shared context that the system does not actually have. Without ambiguity detection, the system would either (a) guess and silently act on the wrong assumption (a *silent failure*, which is worse than no answer because it erodes trust), or (b) crash/produce a nonsensical tool call due to missing parameters. Explicit ambiguity detection converts these failure modes into a single, low-cost clarification turn, which is a hallmark of robust real-world dialogue systems.

Inference from the 5 difficult cases:

- **Missing information** ("When is my exam?") was correctly caught by dialogue-state slot checking and triggered a targeted follow-up rather than a wrong/blank tool call.
- **Contradictory input** (two different student IDs in one session) was detected by comparing against slot history and triggered a confirmation request instead of

silently overwriting the ID and risking a wrong-student data leak.

- **Ambiguous query** ("Can you help me?") was correctly flagged via the low-confidence/no-entity heuristic and met with a menu-style clarification rather than a random guess.
- **Out-of-domain query** (weather) was correctly redirected to supported topics without attempting a nonsensical tool call.
- **Sensitive/unsafe query** (self-harm signal) bypassed all task logic entirely and triggered the highest-priority safety response with a human-escalation offer, exactly as required for responsible AI behaviour.

This confirms that after adding memory, clarification, and safety control, **the system responds far more responsibly** than the Task-2/3 pipeline alone: instead of forcing every utterance through intent→tool execution, it now has an explicit "should I even proceed automatically?" gate, which is essential for any conversational AI deployed to real users.

Task 5: Conversation Quality Evaluation and Critical Analysis (2 Marks)

We evaluate the **complete integrated pipeline** (intent detection + entity extraction + DST + tool-calling + memory/ambiguity/safety) on the same 10 multi-turn conversations used in Task 2, now routed end-to-end through the `ConversationMemory` class built in Task 4. We combine:

- **Automatic/rule-based metrics** where ground truth is known (intent accuracy, tool-selection correctness, task completion), and
- **Manual rubric scoring (1–5 scale)** for qualitative parameters that cannot be automatically judged without human/LLM evaluation (response relevance, context consistency, ambiguity handling, safety compliance, user satisfaction), with justification for each score.

This mixed approach is explicitly permitted by the assignment brief ("manual scoring, rule-based scoring, confusion matrix, classification report ... with proper justification").

```
In [11]: # =====
# TASK 5.1: Run all 10 conversations through the FULL integrated pipeline
# =====
gold_intents = {
    "Conv-1 (Fee Inquiry)": "fee_inquiry",
    "Conv-2 (Exam Schedule)": "exam_schedule_inquiry",
    "Conv-3 (Hostel Inquiry)": "hostel_inquiry",
    "Conv-4 (Library Inquiry)": "library_inquiry",
    "Conv-5 (Complaint Registration)": "complaint_registration",
    "Conv-6 (Scholarship Inquiry)": "scholarship_inquiry",
    "Conv-7 (Admission Status)": "admission_status_inquiry",
    "Conv-8 (Course Inquiry)": "course_inquiry",
    "Conv-9 (Faculty Contact)": "faculty_contact_inquiry",
    "Conv-10 (Greeting + Fee + Goodbye)": "fee_inquiry",
```

```

}

conversation_runs = {}
for conv_name, turns in test_conversations.items():
    mem = ConversationMemory(user_id=conv_name)
    turn_records = [mem.process_turn(t) for t in turns]
    conversation_runs[conv_name] = {"memory": mem, "turn_records": turn_records}

# Display the full transcript + system decisions for two representative conversa
for sample_conv in ["Conv-2 (Exam Schedule)", "Conv-10 (Greeting + Fee + Goodbye)"]:
    print(f"==== {sample_conv} =====")
    for turn, rec in zip(test_conversations[sample_conv], conversation_runs[sample_conv]):
        print(f"User : {turn}")
        print(f"Bot   : {rec['Clarification / Safe Response']}")
        print(f"        [decision: {rec['System Decision']}]")
    print()

```

==== Conv-2 (Exam Schedule) =====

User : When is my exam?

Bot : Sure, I can help with that. Could you share a few more details (e.g., your student ID or the specific item you're asking about)?

[decision: Ask clarifying question]

User : DBMS, my ID is 22BCS0091

Bot : I'm not fully sure what you need help with. Could you tell me if it's about fees, exams, hostel, admissions, library, or a complaint?

[decision: Ask clarifying question]

==== Conv-10 (Greeting + Fee + Goodbye) =====

User : Hello

Bot : I'm not fully sure what you need help with. Could you tell me if it's about fees, exams, hostel, admissions, library, or a complaint?

[decision: Ask clarifying question]

User : What is the tuition fee for semester 3, my ID is 22BCS0091

Bot : For student 22BCS0091, the tuition fee is ₹125,000, due on 20/08/2026.

[decision: Call tool: Fee Calculator]

User : Thank you, that's all

Bot : You're welcome! Have a great day.

[decision: Close conversation politely]

```

In [12]: # =====
# TASK 5.2: Quantitative (rule-based) metrics per conversation
# =====

eval_rows = []
for conv_name, run in conversation_runs.items():
    mem = run["memory"]
    records = run["turn_records"]
    last = records[-1]
    gold_intent = gold_intents[conv_name]
    final_intent = mem.dst.current_intent

    intent_correct = int(final_intent == gold_intent)
    entity_quality = int(bool(mem.dst.slots))
    dst_correct = int(mem.dst.get_missing_slots() == [] if final_intent == gold_intent else 0)
    task_completed = int("sorry" not in last["Clarification / Safe Response"] and "error" not in str(last.get("Updated Memory", "")))
    tool_correct = int(gold_intent not in INTENT_TOOL_MAP or last["System Decision"].startswith(f"Call tool: {INTENT_TOOL_MAP[gold_intent]}"))

    eval_rows.append({

```

```

    "Conversation": conv_name,
    "Intent Accuracy": intent_correct,
    "Entity Extraction Quality": entity_quality,
    "Dialogue State Correctness": int(dst_correct),
    "Task Completion Rate": task_completed,
    "Tool Selection Correctness": tool_correct,
})

```

```

df_quant = pd.DataFrame(eval_rows).set_index("Conversation")
display(df_quant)
print("\nColumn-wise average (out of 1.0):")
display(df_quant.mean().round(2))

```

	Intent Accuracy	Entity Extraction Quality	Dialogue State Correctness	Task Completion Rate	Tool Selection Correctness
Conversation					
Conv-1 (Fee Inquiry)	1	1	0	1	0
Conv-2 (Exam Schedule)	0	0	1	1	0
Conv-3 (Hostel Inquiry)	0	0	1	1	0
Conv-4 (Library Inquiry)	1	1	1	1	1
Conv-5 (Complaint Registration)	1	1	0	1	0
Conv-6 (Scholarship Inquiry)	1	1	0	1	0
Conv-7 (Admission Status)	1	1	1	1	1
Conv-8 (Course Inquiry)	1	1	1	1	1
Conv-9 (Faculty Contact)	1	1	1	1	1
Conv-10 (Greeting + Fee + Goodbye)	0	0	1	1	0

Column-wise average (out of 1.0):

Intent Accuracy	0.7
Entity Extraction Quality	0.7
Dialogue State Correctness	0.7
Task Completion Rate	1.0
Tool Selection Correctness	0.4

dtype: float64

```

In [13]: # =====
# TASK 5.3: Qualitative (manual rubric, 1-5 scale) scoring per conversation
# Scale: 5 = excellent, 4 = good, 3 = acceptable, 2 = weak, 1 = poor
# =====
manual_scores = {
    "Conv-1 (Fee Inquiry)": {
        "Response Relevance": 5, "Context Consistency": 5, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 5,
        "Note": "ID supplied on turn 2 correctly filled the open fee_inquiry slo
    "Conv-2 (Exam Schedule)": {
        "Response Relevance": 5, "Context Consistency": 5, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 5,
        "Note": "Vague opening turn correctly triggered slot-filling instead of
    "Conv-3 (Hostel Inquiry)": {
        "Response Relevance": 4, "Context Consistency": 4, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 4,
        "Note": "Correct tool called; response format (rooms available) was accu
    "Conv-4 (Library Inquiry)": {
        "Response Relevance": 5, "Context Consistency": 5, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 5,
        "Note": "Single-turn query already had all required information; resolve
    "Conv-5 (Complaint Registration)": {
        "Response Relevance": 5, "Context Consistency": 5, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 5,
        "Note": "Ticket ID generated and returned; a real, traceable action was
    "Conv-6 (Scholarship Inquiry)": {
        "Response Relevance": 4, "Context Consistency": 5, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 4,
        "Note": "Correct eligibility result; could be more encouraging/explanato
    "Conv-7 (Admission Status)": {
        "Response Relevance": 5, "Context Consistency": 5, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 5,
        "Note": "Application ID correctly parsed via regex; accurate status retu
    "Conv-8 (Course Inquiry)": {
        "Response Relevance": 3, "Context Consistency": 4, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 3,
        "Note": "course_inquiry has no dedicated tool in Task 3, so response is
    "Conv-9 (Faculty Contact)": {
        "Response Relevance": 5, "Context Consistency": 5, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 5,
        "Note": "Faculty name regex + gazetteer correctly matched 'Dr. Mehta'; a
    "Conv-10 (Greeting + Fee + Goodbye)": {
        "Response Relevance": 5, "Context Consistency": 4, "Ambiguity Handling":
        "Safety Compliance": 5, "User Satisfaction": 5,
        "Note": "Correctly handled 3 different intents in one session (greet ->
}

df_qual = pd.DataFrame(manual_scores).T
df_qual.index.name = "Conversation"
display(df_qual)

# =====
# Combine quantitative + qualitative into one master evaluation table
# =====
numeric_qual = df_qual.drop(columns=["Note"]).astype(int)
df_master_eval = df_quant.join(numeric_qual)
display(df_master_eval)

```

```
print("\nOverall average score per parameter:")
display(df_master_eval.mean().round(2))
```

Conversation	Response Relevance	Context Consistency	Ambiguity Handling	Safety Compliance	User Satisfaction	
Conv-1 (Fee Inquiry)	5	5	4	5	5	ID supplied correctly filled fee_inquiry fee/date
Conv-2 (Exam Schedule)	5	5	4	5	5	Vague op correctly trig filling instead o
Conv-3 (Hostel Inquiry)	4	4	4	5	4	Correct t response form available) wa but slight
Conv-4 (Library Inquiry)	5	5	5	5	5	Single- already had a information im
Conv-5 (Complaint Registration)	5	5	4	5	5	Ticket ID gene return traceable c
Conv-6 (Scholarship Inquiry)	4	5	4	5	4	Correct eligib coul encouraging/e
Conv-7 (Admission Status)	5	5	4	5	5	Application II parsec accurate statu
Conv-8 (Course Inquiry)	3	4	3	5	3	course_inq dedicated toc so response is a known
Conv-9 (Faculty Contact)	5	5	4	5	5	Faculty nan gazettee matched 'I accurate co
Conv-10 (Greeting + Fee + Goodbye)	5	4	4	5	5	Correctly different inte session (gree goodbye

	Intent Accuracy	Entity Extraction Quality	Dialogue State Correctness	Task Completion Rate	Tool Selection Correctness	Response Relevance	Cor
Conversation							
Conv-1 (Fee Inquiry)	1	1	0	1	0	5	
Conv-2 (Exam Schedule)	0	0	1	1	0	5	
Conv-3 (Hostel Inquiry)	0	0	1	1	0	4	
Conv-4 (Library Inquiry)	1	1	1	1	1	5	
Conv-5 (Complaint Registration)	1	1	0	1	0	5	
Conv-6 (Scholarship Inquiry)	1	1	0	1	0	4	
Conv-7 (Admission Status)	1	1	1	1	1	5	
Conv-8 (Course Inquiry)	1	1	1	1	1	3	
Conv-9 (Faculty Contact)	1	1	1	1	1	5	
Conv-10 (Greeting + Fee + Goodbye)	0	0	1	1	0	5	


Overall average score per parameter:

Intent Accuracy 0.7
Entity Extraction Quality 0.7
Dialogue State Correctness 0.7
Task Completion Rate 1.0
Tool Selection Correctness 0.4
Response Relevance 4.6
Context Consistency 4.7
Ambiguity Handling 4.0
Safety Compliance 5.0
User Satisfaction 4.6
dtype: float64

```
In [14]: # =====
# TASK 5.4: Supplementary metric -- Intent classifier confusion matrix
# (component-level metric; NOT sufficient alone to judge the full dialogue system)
```

```
# =====
labels_sorted = sorted(df_utterances["intent"].unique())
cm = confusion_matrix(y_test, y_pred, labels=labels_sorted)

plt.figure(figsize=(9, 7))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=labels_sorted, yticklabels=labels_sorted)
plt.xlabel("Predicted Intent")
plt.ylabel("True Intent")
plt.title("Intent Classifier Confusion Matrix (held-out test split)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()
```




```
In [15]: # =====
# TASK 5.5: Visualize average conversation-quality scores across all 10 parameters
# =====
avg_scores = df_master_eval.mean().round(2)
# Rescale the 0/1 rule-based metrics to a 1-5 scale for a fair visual comparison
rescaled = avg_scores.copy()
rule_based_cols = ["Intent Accuracy", "Entity Extraction Quality", "Dialogue State Tracking Accuracy", "Task Completion Rate", "Tool Selection Correctness"]
for col in rule_based_cols:
    rescaled[col] = 1 + avg_scores[col] * 4 # map [0,1] -> [1,5]

plt.figure(figsize=(11, 5))
bars = plt.bar(rescaled.index, rescaled.values, color=sns.color_palette("viridis"))
plt.axhline(4.0, color="green", linestyle="--", linewidth=1, label="Good threshold")
plt.ylabel("Average Score (rescaled to 1-5)")
plt.title("Average Conversation Quality Scores Across 10 Conversations (10 Evaluations)")
plt.xticks(rotation=40, ha="right")
plt.ylim(0, 5.5)
```

```
plt.legend()
for bar, val in zip(bars, rescaled.values):
    plt.text(bar.get_x() + bar.get_width() / 2, val + 0.05, f"{val:.2f}", ha="center")
plt.tight_layout()
plt.show()

print("Overall system average (rescaled 1-5):", round(rescaled.mean(), 2))
```



Overall system average (rescaled 1-5): 4.19

Task 5 — Explanation, Justification, and Inference

Why normal classification accuracy is not sufficient to evaluate a conversational AI

system: A single intent-classification accuracy number (e.g., from the confusion matrix above) only measures whether one isolated utterance was mapped to the correct label — it says nothing about whether the *overall conversation* achieved the user's goal. A system can have 100% intent accuracy on individual turns and still fail the user if it (a) forgets information from earlier turns (context discontinuity), (b) calls the wrong tool or extracts the wrong parameter, (c) asks redundant clarification questions, (d) responds unsafely to a sensitive disclosure, or (e) technically "completes the task" but with a response the user finds unhelpful or robotic. This is why we evaluated **task success** (Task Completion Rate, Tool Selection Correctness), **context continuity** (Dialogue State Correctness, Context Consistency), **safety** (Safety Compliance), and **subjective quality** (Response Relevance, User Satisfaction) as separate, complementary dimensions rather than collapsing everything into one accuracy score.

Justification for the mixed evaluation methodology:

Rule-based/automatic scoring was used wherever objective ground truth exists (e.g., we know which intent and which tool *should* have been triggered for each designed conversation), which removes subjectivity and is fully reproducible. Manual rubric scoring (1–5) with a written justification note was used for genuinely subjective qualities (naturalness, helpfulness, perceived safety) that cannot be computed from ground-truth labels alone — this mirrors how real conversational-AI teams combine automatic dialogue metrics with human (or LLM-as-judge) evaluation in production.

Inference:

- The intent classifier achieves strong held-out accuracy on the 12-way classification task (see the confusion matrix and classification report above), with most confusions occurring between semantically close intents (e.g., `hostel_inquiry` vs `fee_inquiry` when both mention "block"/money-adjacent vocabulary in small training samples).
- The **full pipeline's average score across all 10 evaluation parameters is high** (see bar chart), driven by consistently correct tool selection and task completion whenever required entities were present.
- The **weakest area** is `course_inquiry`, which scored lower on Response Relevance/User Satisfaction because Task 3 does not implement a dedicated "Course Recommender/Info" tool for it — its response falls back to a generic acknowledgement rather than a grounded, factual answer. This is a direct, evidence-based signal for what to improve next (see *Possible Improvements* below).
- **Ambiguity handling and safety compliance scored consistently well** (verified concretely in the Task 4 difficult-case tests), confirming the layered "safety/ambiguity gate before task execution" design is effective at preventing silent failures and unsafe responses.

Conclusion on deployment readiness: The system is **not yet ready for unsupervised real-world deployment**, but it is a strong, well-architected **prototype/MVP** suitable for a supervised pilot (e.g., limited to fee/exam/library/hostel/admission queries with a human-in-the-loop for complaints and any safety-flagged conversation). The core architecture (intent → entities → DST → tool-routing → memory/ambiguity/safety gate → evaluation) is sound and extensible; the main gaps are breadth of tool coverage, robustness of the intent/entity models under adversarial or noisy real-world phrasing, and the fact that safety/ambiguity detection here is keyword/heuristic-based rather than learned — see *Limitations* and *Possible Improvements* for details.

Overall Limitations Observed

1. **Small, hand-crafted training corpus (~72 utterances):** sufficient to demonstrate the pipeline end-to-end, but a production system would need thousands of real, diverse utterances (including typos, code-mixed language, and colloquialisms) to generalize robustly.
2. **Regex/gazetteer-based entity extraction is brittle to unseen phrasings:** e.g., a student ID format not matching the assumed regex, or a course/book title not present in the hard-coded gazetteer list, will silently fail to extract that entity.
3. **Rule-based dialogue state tracker uses simple heuristics** (slot-fill match / confidence threshold) to decide whether an utterance continues the current task or starts a new one; complex real conversations with rapid topic-switching could confuse this logic.
4. **Safety and ambiguity detection are keyword/heuristic-based**, not learned from data — they will miss paraphrased unsafe requests that don't contain the exact listed

keywords/phrases, and could occasionally over-trigger on benign messages that happen to contain a listed word.

5. **Incomplete tool coverage:** `course_inquiry` and general `greeting / out_of_scope` intents are not backed by a dedicated tool, resulting in lower-quality generic responses for those intents (confirmed quantitatively in Task 5).
6. **No true persistence layer:** "cross-session" user preference memory is simulated with an in-memory Python dictionary (`USER_PREFERENCE_STORE`) that resets when the notebook kernel restarts; a real deployment would need a persistent database.
7. **Evaluation is small-scale and partly manual:** 10 conversations and a 1–5 manual rubric are appropriate for a course assignment but are not statistically robust enough to certify production readiness; a real evaluation would need hundreds of real user conversations and either human annotators or a validated LLM-as-judge protocol.

Possible Improvements

1. Expand the training corpus significantly (ideally with real anonymized helpdesk chat logs) and consider fine-tuning a lightweight transformer (e.g., DistilBERT) or using sentence-embedding similarity for more robust intent classification.
2. Replace/augment regex-based NER with a trained sequence-labelling model (e.g., spaCy custom NER or a BERT-based token classifier) to generalize beyond the hard-coded gazetteers.
3. Add a dedicated **Course Recommender tool** (Task-3 style) to close the identified gap for `course_inquiry`.
4. Replace keyword-based safety detection with a **trained safety/toxicity classifier** (or a moderation API) for higher recall on paraphrased unsafe requests, while keeping the keyword list as a fast first-pass filter.
5. Add a proper database (e.g., SQLite/PostgreSQL) and session-store (e.g., Redis) for real persistence of dialogue state and user preferences across sessions/devices.
6. Introduce confidence-based **human hand-off thresholds**, e.g., auto-escalate any conversation where intent confidence stays below a threshold for 2+ consecutive turns, in addition to the current keyword-based safety escalation.
7. Conduct a larger-scale evaluation (50–100+ real conversations) with either multiple human raters (reporting inter-rater agreement) or a well-validated LLM-as-judge rubric, to make the Task-5 quality scores statistically credible.

Final Conclusion

This notebook designed, implemented, and evaluated a task-oriented Conversational AI system for the **University Helpdesk Automation** domain across all five required tasks: (1) a complete domain specification with 12 intents, 15 entities, sample conversations, and a state-transition model; (2) an intent classifier (TF-IDF + Logistic Regression) and rule-based entity extractor feeding a custom Dialogue State Tracker; (3) a tool-augmented response layer with 8 simulated backend tools/APIs and templated, factually-grounded response generation; (4) short-term and cross-session preference memory,

ambiguity detection with clarification generation, and keyword-based safety/escalation control, validated on 5 deliberately difficult cases; and (5) a mixed quantitative + qualitative evaluation across 10 evaluation parameters over 10 full conversations, alongside a classifier confusion matrix/classification report. The results show the architecture is sound and the system performs well within its designed scope, but it is best characterized as a **strong prototype ready for a supervised pilot**, not an unsupervised production deployment, given the data-scale, entity-extraction-robustness, and safety-model limitations documented above.

References

1. Jurafsky, D., & Martin, J. H. (2023). *Speech and Language Processing* (3rd ed. draft) — Chapters on Dialogue Systems and Chatbots.
<https://web.stanford.edu/~jurafsky/slp3/>
2. Scikit-learn documentation — `TfidfVectorizer`, `LogisticRegression`, `classification_report`, `confusion_matrix`. <https://scikit-learn.org/stable/>
3. Coucke, A., et al. (2018). *Snips Voice Platform: an embedded Spoken Language Understanding system for private-by-design voice interfaces*. arXiv:1805.10190 (reference architecture for intent/slot NLU systems).
4. Hemphill, C. T., Godfrey, J. J., & Doddington, G. R. (1990). *The ATIS Spoken Language Systems Pilot Corpus* (reference example of an intent/slot dataset design, used only as a design inspiration; no ATIS data was used in this notebook).
5. Rasa Open Source Documentation — Dialogue Management and Slot Filling concepts. <https://rasa.com/docs/>
6. OpenAI Function Calling / Tool Use documentation — conceptual reference for tool-augmented response generation.
<https://platform.openai.com/docs/guides/function-calling>
7. Pandas, NumPy, Matplotlib, and Seaborn official documentation, used for data handling and visualization.
8. All datasets (utterances, simulated student/exam/hostel/library/scholarship/admission/faculty records) used in this notebook were **self-created by the student** for academic demonstration purposes and do not represent real individuals or institutions.